

✓ Starter Notebook

Install and import required libraries

```
!pip install transformers datasets evaluate accelerate peft trl bitsandbytes
!pip install nvidia-ml-py3
!pip install scikit-learn
```

```
⇒ Requirement already satisfied: safetensors<0.4.4,>=0.4.3 in /home/ab12660/.local/lib/python3.9/site-packages/safetensors-0.4.3-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: tqdm<4.67,>=4.27 in /home/ab12660/.local/lib/python3.9/site-packages/tqdm-4.66.1-py3-none-any.whl
Requirement already satisfied: pyarrow<16.0.0,>=15.0.0 in /home/ab12660/.local/lib/python3.9/site-packages/pyarrow-15.0.2-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: dill<0.3.9,>=0.3.0 in /home/ab12660/.local/lib/python3.9/site-packages/dill-0.3.8-py3-none-any.whl
Requirement already satisfied: pandas in /home/ab12660/.local/lib/python3.9/site-packages/pandas-2.1.4-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: xxhash in /home/ab12660/.local/lib/python3.9/site-packages/xxhash-3.4.1-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: multiprocessing<0.71,>=0.70.17 in /home/ab12660/.local/lib/python3.9/site-packages/multiprocessing-0.70.17-py3-none-any.whl
Requirement already satisfied: fsspec<2025.0.0,>=2024.12.0 in /home/ab12660/.local/lib/python3.9/site-packages/fsspec-2024.12.0-py3-none-any.whl
Requirement already satisfied: aiohttp in /home/ab12660/.local/lib/python3.9/site-packages/aiohttp-3.9.5-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: psutil in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/psutil-5.9.8-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: torch<2.5.0,>=2.0.0 in /home/ab12660/.local/lib/python3.9/site-packages/torch-2.4.0-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: rich in /home/ab12660/.local/lib/python3.9/site-packages/rich-13.7.1-py3-none-any.whl
Requirement already satisfied: aiohappyeyeballs<2.4.0,>=2.3.0 in /home/ab12660/.local/lib/python3.9/site-packages/aiohappyeyeballs-2.3.5-py3-none-any.whl
Requirement already satisfied: aiosignal<2.0.0,>=1.1.2 in /home/ab12660/.local/lib/python3.9/site-packages/aiosignal-1.1.2-py3-none-any.whl
Requirement already satisfied: async-timeout<6.0,>=4.0 in /home/ab12660/.local/lib/python3.9/site-packages/async_timeout-4.0.3-py3-none-any.whl
Requirement already satisfied: attrs<25.0.0,>=17.3.0 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/attrs-24.2.0-py3-none-any.whl
Requirement already satisfied: frozenlist<1.4.0,>=1.1.1 in /home/ab12660/.local/lib/python3.9/site-packages/frozenlist-1.4.0-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: multidict<7.0,>=4.5 in /home/ab12660/.local/lib/python3.9/site-packages/multidict-6.0.5-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: propcache<0.2.0,>=0.2.0 in /home/ab12660/.local/lib/python3.9/site-packages/propcache-0.2.0-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: yarl<2.0,>=1.17.0 in /home/ab12660/.local/lib/python3.9/site-packages/yarl-1.17.0-cp39-cp39-linux_x86_64.whl
Requirement already satisfied: typing-extensions<4.12.0,>=3.7.4.3 in /home/ab12660/.local/lib/python3.9/site-packages/typing_extensions-4.11.0-py3-none-any.whl
Requirement already satisfied: charset-normalizer<4,>=2 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/charset-normalizer-3.3.0-py3-none-any.whl
Requirement already satisfied: idna<4,>=2.5 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/idna-3.10-py3-none-any.whl
Requirement already satisfied: urllib3<3,>=1.21.1 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/urllib3-2.2.3-py3-none-any.whl
Requirement already satisfied: certifi<2025.0.0,>=2017.4.17 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/certifi-2024.12.14-py3-none-any.whl
Requirement already satisfied: networkx in /home/ab12660/.local/lib/python3.9/site-packages/networkx-3.3-py3-none-any.whl
Requirement already satisfied: jinja2 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/jinja2-3.1.3-py3-none-any.whl
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.4.127 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-any.whl
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.4.127 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_cuda_runtime_cu12-12.4.127-py3-none-any.whl
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.4.127 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_cuda_cupti_cu12-12.4.127-py3-none-any.whl
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_cudnn_cu12-9.1.0.70-py3-none-any.whl
Requirement already satisfied: nvidia-cublas-cu12==12.4.5.8 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_cublas_cu12-12.4.5.8-py3-none-any.whl
Requirement already satisfied: nvidia-cufft-cu12==11.2.1.3 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_cufft_cu12-11.2.1.3-py3-none-any.whl
Requirement already satisfied: nvidia-curand-cu12==10.3.5.147 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_curand_cu12-10.3.5.147-py3-none-any.whl
Requirement already satisfied: nvidia-cusolver-cu12==11.6.1.9 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_cusolver_cu12-11.6.1.9-py3-none-any.whl
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_cusparselt_cu12-0.6.2-py3-none-any.whl
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_nccl_cu12-2.21.5-py3-none-any.whl
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /home/ab12660/.local/lib/python3.9/site-packages/nvidia_nvtx_cu12-12.4.127-py3-none-any.whl
```

```

Requirement already satisfied: nvidia-nvjitlink-cu12==12.4.127 in /home/ab12
Requirement already satisfied: triton==3.2.0 in /home/ab12660/.local/lib/pyt
Requirement already satisfied: sympy==1.13.1 in /home/ab12660/.local/lib/pyt
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /home/ab12660/.local/li
Requirement already satisfied: python-dateutil>=2.8.2 in /share/apps/pyenv/p
Requirement already satisfied: pytz>=2020.1 in /home/ab12660/.local/lib/pyth
Requirement already satisfied: tzdata>=2022.7 in /home/ab12660/.local/lib/py
Requirement already satisfied: markdown-it-py>=2.2.0 in /home/ab12660/.local
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /share/apps/pyenv/
Requirement already satisfied: mdurl~=0.1 in /home/ab12660/.local/lib/pythor
Requirement already satisfied: six>=1.5 in /share/apps/pyenv/py3.9/lib/pythc
Requirement already satisfied: MarkupSafe>=2.0 in /share/apps/pyenv/py3.9/li
Defaulting to user installation because normal site-packages is not writeabl
Requirement already satisfied: nvidia-ml-py3 in /home/ab12660/.local/lib/pyt
Defaulting to user installation because normal site-packages is not writeabl
Requirement already satisfied: scikit-learn in /home/ab12660/.local/lib/pyth
Requirement already satisfied: numpy>=1.19.5 in /home/ab12660/.local/lib/pyt
Requirement already satisfied: scipy>=1.6.0 in /home/ab12660/.local/lib/pyth
Requirement already satisfied: joblib>=1.2.0 in /home/ab12660/.local/lib/pyt
Requirement already satisfied: threadpoolctl>=3.1.0 in /home/ab12660/.local/

```

```

!pip install os
!pip install pandas
!pip install torch
!pip install transformers
!pip install peft
!pip install datasets
!pip install pickle
!pip install --user pandas

```

```

⇒ Defaulting to user installation because normal site-packages is not writeabl
ERROR: Could not find a version that satisfies the requirement os (from vers
ERROR: No matching distribution found for os
Defaulting to user installation because normal site-packages is not writeabl
Requirement already satisfied: pandas in /home/ab12660/.local/lib/python3.9/
Requirement already satisfied: numpy>=1.22.4 in /home/ab12660/.local/lib/pyt
Requirement already satisfied: python-dateutil>=2.8.2 in /share/apps/pyenv/p
Requirement already satisfied: pytz>=2020.1 in /home/ab12660/.local/lib/pyth
Requirement already satisfied: tzdata>=2022.7 in /home/ab12660/.local/lib/py
Requirement already satisfied: six>=1.5 in /share/apps/pyenv/py3.9/lib/pythc
Defaulting to user installation because normal site-packages is not writeabl
Requirement already satisfied: torch in /home/ab12660/.local/lib/python3.9/s
Requirement already satisfied: filelock in /home/ab12660/.local/lib/python3.
Requirement already satisfied: typing-extensions>=4.10.0 in /home/ab12660/.l
Requirement already satisfied: networkx in /home/ab12660/.local/lib/python3.
Requirement already satisfied: jinja2 in /share/apps/pyenv/py3.9/lib/python3
Requirement already satisfied: fsspec in /home/ab12660/.local/lib/python3.9/
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.4.127 in /home/ab1
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.4.127 in /home/a
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.4.127 in /home/ab1
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in /home/ab12660/

```

Requirement already satisfied: nvidia-cublas-cu12==12.4.5.8 in /home/ab12660/

Requirement already satisfied: nvidia-cufft-cu12==11.2.1.3 in /home/ab12660/

Requirement already satisfied: nvidia-curand-cu12==10.3.5.147 in /home/ab12660/

Requirement already satisfied: nvidia-cusolver-cu12==11.6.1.9 in /home/ab12660/

Requirement already satisfied: nvidia-cusparselt-cu12==12.3.1.170 in /home/ab12660/

Requirement already satisfied: nvidia-cusparse-cu12==12.3.1.170 in /home/ab12660/

Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /home/ab12660/

Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /home/ab12660/.local/lib/python3.9/

Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /home/ab12660/.local/lib/python3.9/

Requirement already satisfied: nvidia-nvjitlink-cu12==12.4.127 in /home/ab12660/.local/lib/python3.9/

Requirement already satisfied: triton==3.2.0 in /home/ab12660/.local/lib/python3.9/

Requirement already satisfied: sympy==1.13.1 in /home/ab12660/.local/lib/python3.9/

Requirement already satisfied: mpmath<1.4,>=1.1.0 in /home/ab12660/.local/lib/python3.9/

Requirement already satisfied: MarkupSafe>=2.0 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Defaulting to user installation because normal site-packages is not writeable

Requirement already satisfied: transformers in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: filelock in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: huggingface-hub<1.0,>=0.30.0 in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: numpy>=1.17 in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: packaging>=20.0 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Requirement already satisfied: pyyaml>=5.1 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Requirement already satisfied: regex!=2019.12.17 in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: requests in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: tokenizers<0.22,>=0.21 in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: safetensors>=0.4.3 in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: tqdm>=4.27 in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: fsspec>=2023.5.0 in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: typing-extensions>=3.7.4.3 in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: charset-normalizer<4,>=2 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Requirement already satisfied: idna<4,>=2.5 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Requirement already satisfied: urllib3<3,>=1.21.1 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Requirement already satisfied: certifi>=2017.4.17 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Defaulting to user installation because normal site-packages is not writeable

Requirement already satisfied: peft in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: numpy>=1.17 in /home/ab12660/.local/lib/python3.9/site-packages/

Requirement already satisfied: packaging>=20.0 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Requirement already satisfied: psutil in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Requirement already satisfied: pyyaml in /share/apps/pyenv/py3.9/lib/python3.9/site-packages/

Requirement already satisfied: torch>=1.13.0 in /home/ab12660/.local/lib/python3.9/site-packages/

```

!pip install pandas
import os
import pandas as pd
import torch
from transformers import RobertaModel, RobertaTokenizer, TrainingArguments, Traini
from peft import LoraConfig, get_peft_model, PeftModel
from datasets import load_dataset, Dataset, ClassLabel
import pickle

```

```

⇒ Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: pandas in /home/ab12660/.local/lib/python3.9/s:
Requirement already satisfied: numpy>=1.22.4 in /home/ab12660/.local/lib/pytho
Requirement already satisfied: python-dateutil>=2.8.2 in /share/apps/pyenv/py:
Requirement already satisfied: pytz>=2020.1 in /home/ab12660/.local/lib/pytho
Requirement already satisfied: tzdata>=2022.7 in /home/ab12660/.local/lib/pytl
Requirement already satisfied: six>=1.5 in /share/apps/pyenv/py3.9/lib/python:
/home/ab12660/.local/lib/python3.9/site-packages/tqdm/auto.py:21: TqdmWarning:
    from .autonotebook import tqdm as notebook_tqdm

```

✓ Load Tokenizer and Preprocess Data

Start coding or [generate](#) with AI.

```
base_model = 'roberta-base'
```

```
dataset = load_dataset('ag_news', split='train')
tokenizer = RobertaTokenizer.from_pretrained(base_model)
```

```
def preprocess(examples):
    tokenized = tokenizer(examples['text'], truncation=True, padding=True)
    return tokenized
```

```
tokenized_dataset = dataset.map(preprocess, batched=True, remove_columns=["text"]
tokenized_dataset = tokenized_dataset.rename_column("label", "labels")
```

```
# Extract the number of classes and their names
num_labels = dataset.features['label'].num_classes
class_names = dataset.features["label"].names
print(f"number of labels: {num_labels}")
print(f"the labels: {class_names}")

# Create an id2label mapping
# We will need this for our classifier.
id2label = {i: label for i, label in enumerate(class_names)}

data_collator = DataCollatorWithPadding(tokenizer=tokenizer, return_tensors="pt")
```

```
➡ number of labels: 4
   the labels: ['World', 'Sports', 'Business', 'Sci/Tech']
```

✓ Load Pre-trained Model

Set up config for pretrained model and download it from hugging face

```
model = RobertaForSequenceClassification.from_pretrained(
    base_model,
    id2label=id2label)
model
```

⇒ Some weights of RobertaForSequenceClassification were not initialized from the pre-trained weights. You should probably TRAIN this model on a down-stream task to be able to use it.

```
RobertaForSequenceClassification(  
  (roberta): RobertaModel(  
    (embeddings): RobertaEmbeddings(  
      (word_embeddings): Embedding(50265, 768, padding_idx=1)  
      (position_embeddings): Embedding(514, 768, padding_idx=1)  
      (token_type_embeddings): Embedding(1, 768)  
      (LayerNorm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)  
      (dropout): Dropout(p=0.1, inplace=False)  
    )  
    (encoder): RobertaEncoder(  
      (layer): ModuleList(  
        (0-11): 12 x RobertaLayer(  
          (attention): RobertaAttention(  
            (self): RobertaSdpaSelfAttention(  
              (query): Linear(in_features=768, out_features=768, bias=True)  
              (key): Linear(in_features=768, out_features=768, bias=True)  
              (value): Linear(in_features=768, out_features=768, bias=True)  
              (dropout): Dropout(p=0.1, inplace=False)  
            )  
            (output): RobertaSelfOutput(  
              (dense): Linear(in_features=768, out_features=768, bias=True)  
              (LayerNorm): LayerNorm((768,), eps=1e-05,  
elementwise_affine=True)  
              (dropout): Dropout(p=0.1, inplace=False)  
            )  
          )  
          (intermediate): RobertaIntermediate(  
            (dense): Linear(in_features=768, out_features=3072, bias=True)  
            (intermediate_act_fn): GELUActivation()  
          )  
          (output): RobertaOutput(  
            (dense): Linear(in_features=3072, out_features=768, bias=True)  
            (LayerNorm): LayerNorm((768,), eps=1e-05,  
elementwise_affine=True)  
            (dropout): Dropout(p=0.1, inplace=False)  
          )  
        )  
      )  
    )  
  )  
  (classifier): RobertaClassificationHead(  
    (dense): Linear(in_features=768, out_features=768, bias=True)  
    (dropout): Dropout(p=0.1, inplace=False)  
    (out_proj): Linear(in_features=768, out_features=4, bias=True)  
  )  
)
```

✓ Anything from here on can be modified

```
# Split the original training set
split_datasets = tokenized_dataset.train_test_split(test_size=640, seed=42)
train_dataset = split_datasets['train']
eval_dataset = split_datasets['test']
```

✓ Setup LoRA Config

Setup PEFT config and get peft model for finetuning

```
# PEFT Config
# Replace old LoRA config with this
peft_config = LoraConfig(
    r=10,                                # Try r=10 for stronger adaptation
    lora_alpha=40,                        # Stronger low-rank scaling
    lora_dropout=0.05,
    bias='none',
    target_modules=["query", "value"], # Include both query and value projection
    task_type="SEQ_CLS"
)
```

```
peft_model = get_peft_model(model, peft_config)
peft_model
```

```
⇒ PeftModelForSequenceClassification(
  (base_model): LoraModel(
    (model): RobertaForSequenceClassification(
      (roberta): RobertaModel(
        (embeddings): RobertaEmbeddings(
          (word_embeddings): Embedding(50265, 768, padding_idx=1)
          (position_embeddings): Embedding(514, 768, padding_idx=1)
          (token_type_embeddings): Embedding(1, 768)
          (LayerNorm): LayerNorm((768,), eps=1e-05,
elementwise_affine=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
        (encoder): RobertaEncoder(
          (layer): ModuleList(
            (0-11): 12 x RobertaLayer(
              (attention): RobertaAttention(
                (self): RobertaSdpaSelfAttention(
                  (query): lora.Linear(
```

```

        (base_layer): Linear(in_features=768, out_features=768,
bias=True)
        (lora_dropout): ModuleDict(
          (default): Dropout(p=0.05, inplace=False)
        )
        (lora_A): ModuleDict(
          (default): Linear(in_features=768, out_features=10,
bias=False)
        )
        (lora_B): ModuleDict(
          (default): Linear(in_features=10, out_features=768,
bias=False)
        )
        (lora_embedding_A): ParameterDict()
        (lora_embedding_B): ParameterDict()
        (lora_magnitude_vector): ModuleDict()
      )
    (key): Linear(in_features=768, out_features=768,
bias=True)
    (value): lora.Linear(
      (base_layer): Linear(in_features=768, out_features=768,
bias=True)
      (lora_dropout): ModuleDict(
        (default): Dropout(p=0.05, inplace=False)
      )
      (lora_A): ModuleDict(
        (default): Linear(in_features=768, out_features=10,
bias=False)
      )
      (lora_B): ModuleDict(
        (default): Linear(in_features=10, out_features=768,
bias=False)
      )
      (lora_embedding_A): ParameterDict()
      (lora_embedding_B): ParameterDict()
      (lora_magnitude_vector): ModuleDict()
    )
    (dropout): Dropout(p=0.1, inplace=False)
  )
  (output): RobertaSelfOutput(
    (dense): Linear(in_features=768, out_features=768,

```

```

print("Trainable parameters:")
for name, param in peft_model.named_parameters():
    if param.requires_grad:
        print(name)

```

➡ Trainable parameters:

```

base_model.model.roberta.encoder.layer.0.attention.self.query.lora_A.default.\
base_model.model.roberta.encoder.layer.0.attention.self.query.lora_B.default.\
base_model.model.roberta.encoder.layer.0.attention.self.value.lora_A.default.\
base_model.model.roberta.encoder.layer.0.attention.self.value.lora_B.default.\
base_model.model.roberta.encoder.layer.1.attention.self.query.lora_A.default.\

```


[illegible]

```
print('PEFT Model')
peft_model.print_trainable_parameters()
```

```
↔ PEFT Model
   trainable params: 962,308 || all params: 125,611,016 || trainable%: 0.7661
```

✓ Training Setup

```
# To track evaluation accuracy during training
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

def compute_metrics(pred):
    labels = pred.label_ids
    preds = pred.predictions.argmax(-1)
    # Calculate accuracy
    accuracy = accuracy_score(labels, preds)
    return {
        'accuracy': accuracy
    }
```

```
# Setup Training args
output_dir = "results"
# Better training setup
from transformers import TrainingArguments
```


```
training_args = TrainingArguments(
    output_dir="results",
    eval_strategy="epoch",
    save_strategy="epoch",
    learning_rate=2e-5,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=64,
    num_train_epochs=6,
    weight_decay=0.01,
    warmup_ratio=0.1,
    lr_scheduler_type="linear",
    load_best_model_at_end=True,
    metric_for_best_model="accuracy",
    greater_is_better=True,
    fp16=True,
    optim="adamw_torch",
    logging_dir="./logs",
    report_to=None
)
```

```
def get_trainer(model):
    return Trainer(
        model=model,
        args=training_args,
        compute_metrics=compute_metrics,
        train_dataset=train_dataset,
        eval_dataset=eval_dataset,
        data_collator=data_collator,
    )
```

✓ Start Training

```
peft_lora_finetuning_trainer = get_trainer(peft_model)
```

```
result = peft_lora_finetuning_trainer.train()
```

 No label_names provided for model class `PeftModelForSequenceClassification`.
[44760/44760 1:05:05, Epoch 6/6]

Epoch	Training Loss	Validation Loss	Accuracy
1	0.265600	0.324273	0.900000
2	0.245200	0.268220	0.914062
3	0.217600	0.242825	0.928125
4	0.208900	0.234309	0.926562
5	0.202100	0.230531	0.931250
6	0.193800	0.236640	0.925000

✓ Evaluate Finetuned Model

Double-click (or enter) to edit

✓ Performing Inference on Custom Input

Uncomment following functions for running inference on custom inputs

```
def classify(model, tokenizer, text):
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    inputs = tokenizer(text, truncation=True, padding=True, return_tensors="pt")
    output = model(**inputs)

    prediction = output.logits.argmax(dim=-1).item()

    print(f'\n Class: {prediction}, Label: {id2label[prediction]}, Text: {text}')
    return id2label[prediction]

classify(peft_model, tokenizer, "Kederis proclaims innocence Olympic champion K")
classify(peft_model, tokenizer, "Wall St. Bears Claw Back Into the Black (Reute
```

⇒

```
Class: 0, Label: World, Text: Kederis proclaims innocence Olympic champion K
Class: 2, Label: Business, Text: Wall St. Bears Claw Back Into the Black (Reu
'Business'
```

✓ Run Inference on eval_dataset

```
from torch.utils.data import DataLoader
import evaluate
from tqdm import tqdm

def evaluate_model(inference_model, dataset, labelled=True, batch_size=8, data_co
    """
    Evaluate a PEFT model on a dataset.

    Args:
        inference_model: The model to evaluate.
        dataset: The dataset (Hugging Face Dataset) to run inference on.
        labelled (bool): If True, the dataset includes labels and metrics will be
                        If False, only predictions will be returned.
        batch_size (int): Batch size for inference.
        data_collator: Function to collate batches. If None, the default collate_

    Returns:
        If labelled is True, returns a tuple (metrics, predictions)
        If labelled is False, returns the predictions.
    """
    # Create the DataLoader
    eval_dataloader = DataLoader(dataset, batch_size=batch_size, collate_fn=data_
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```

inference_model.to(device)
inference_model.eval()

all_predictions = []
if labelled:
    metric = evaluate.load('accuracy')

# Loop over the DataLoader
for batch in tqdm(eval_dataloader):
    # Move each tensor in the batch to the device
    batch = {k: v.to(device) for k, v in batch.items()}
    with torch.no_grad():
        outputs = inference_model(**batch)
        predictions = outputs.logits.argmax(dim=-1)
        all_predictions.append(predictions.cpu())

    if labelled:
        # Expecting that labels are provided under the "labels" key.
        references = batch["labels"]
        metric.add_batch(
            predictions=predictions.cpu().numpy(),
            references=references.cpu().numpy()
        )

# Concatenate predictions from all batches
all_predictions = torch.cat(all_predictions, dim=0)

if labelled:
    eval_metric = metric.compute()
    print("Evaluation Metric:", eval_metric)
    return eval_metric, all_predictions
else:
    return all_predictions

```

```

# Check evaluation accuracy

```

```

_, _ = evaluate_model(peft_model, eval_dataset, True, 8, data_collator)

```

```

➡ 100%|██████████| 80/80 [00:01<00:00, 49.58it/s]Evaluation Metric: {'accuracy':

```

```
# Print final evaluation accuracy
final_eval = peft_lora_finetuning_trainer.evaluate()
print(f"\n✅ Final Eval Accuracy: {final_eval['eval_accuracy']:.4f}")

# Show number of parameters
peft_model.print_trainable_parameters()
```

➡ [10/10 00:01]

✅ Final Eval Accuracy: 0.9313
 trainable params: 962,308 || all params: 125,611,016 || trainable%: 0.7661

Start coding or [generate](#) with AI.

✓ Run Inference on unlabelled dataset

```
#Load your unlabelled data
unlabelled_dataset = pd.read_pickle("test_unlabelled.pkl")
test_dataset = unlabelled_dataset.map(preprocess, batched=True, remove_columns=["·
unlabelled_dataset
```

➡ Map: 100%|██████████| 8000/8000 [00:05<00:00, 1429.69 examples/s]
 Dataset({
 features: ['text'],
 num_rows: 8000
 })

```
# Run inference and save predictions
preds = evaluate_model(peft_model, test_dataset, False, 8, data_collator)
df_output = pd.DataFrame({
    'ID': range(len(preds)),
    'Label': preds.numpy() # or preds.tolist()
})
df_output.to_csv(os.path.join(output_dir,"inference_output.csv"), index=False)
print("Inference complete. Predictions saved to inference_output.csv")
```

➡ 100%|██████████| 1000/1000 [00:18<00:00, 54.10it/s]
 Inference complete. Predictions saved to inference_output.csv

```
# Save model and tokenizer for reproducibility
peft_model.save_pretrained("./results/final_lora_model")
tokenizer.save_pretrained("./results/final_lora_model")
```

```
↔ ('./results/final_lora_model/tokenizer_config.json',
   './results/final_lora_model/special_tokens_map.json',
   './results/final_lora_model/vocab.json',
   './results/final_lora_model/merges.txt',
   './results/final_lora_model/added_tokens.json')
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
def print_metrics(y_true, y_pred, dataset_name=""):
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred, average='macro', zero_division=0)
    rec = recall_score(y_true, y_pred, average='macro', zero_division=0)
    f1 = f1_score(y_true, y_pred, average='macro', zero_division=0)

    print(f"\n=== {dataset_name} Metrics ===")
    print(classification_report(y_true, y_pred, target_names=class_names, zero_division=0))
    return acc, prec, rec, f1
```



```

# Utility function to extract predictions and labels from a dataset
from torch.utils.data import DataLoader

def get_predictions(model, dataset, batch_size=8, data_collator=None):
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    dataloader = DataLoader(dataset, batch_size=batch_size, collate_fn=data_collator)

    model.to(device)
    model.eval()

    all_preds = []
    all_labels = []

    for batch in dataloader:
        labels = batch["labels"].to(device)
        inputs = {k: v.to(device) for k, v in batch.items() if k != "labels"}
        with torch.no_grad():
            outputs = model(**inputs)
            preds = outputs.logits.argmax(dim=-1)

        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

    return all_labels, all_preds

train_labels, train_preds = get_predictions(peft_model, train_dataset, data_collator=train_collator)
test_labels, test_preds = get_predictions(peft_model, eval_dataset, data_collator=test_collator)

# Run this code after getting predictions to generate plots

from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Accuracy comparison
train_acc = accuracy_score(train_labels, train_preds)
test_acc = accuracy_score(test_labels, test_preds)

plt.figure(figsize=(6, 4))
plt.bar(['Train Accuracy', 'Test Accuracy'], [train_acc, test_acc], color=['skyblue', 'lightcoral'])
plt.ylim(0, 1)
plt.ylabel('Accuracy')

```


```


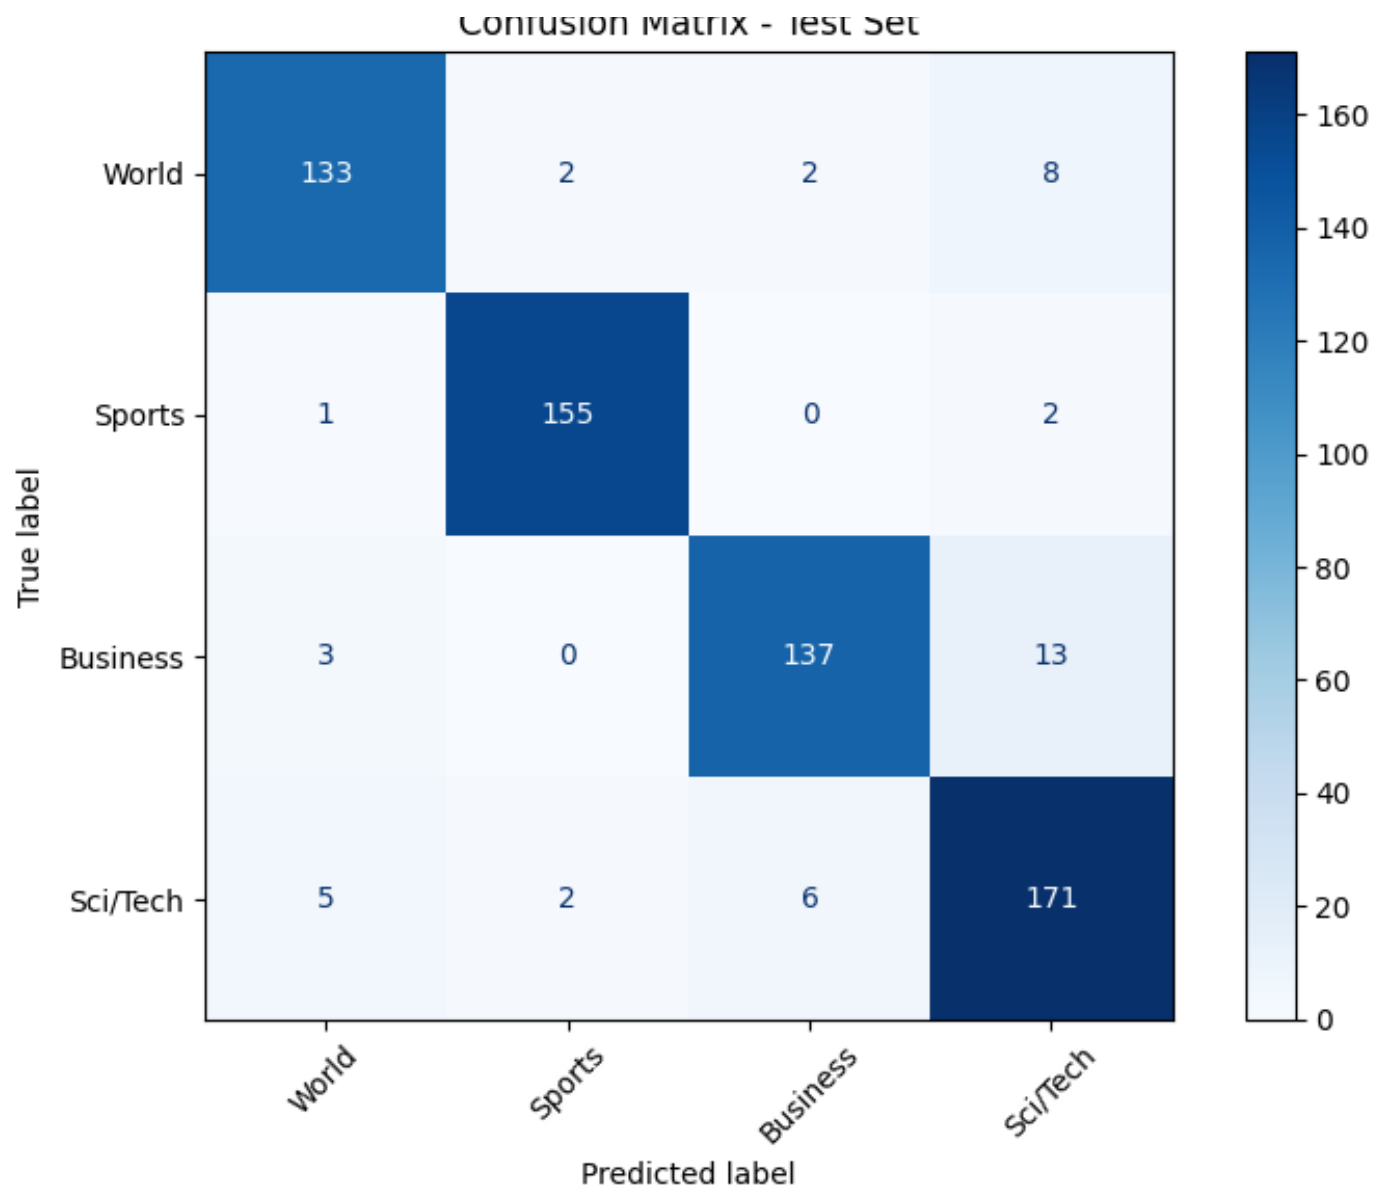
plt.title('Train vs Test Accuracy')
plt.grid(axis='y')
plt.tight_layout()
plt.savefig("train_vs_test_accuracy.png")
plt.show()

# Confusion matrix
cm = confusion_matrix(test_labels, test_preds)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
fig, ax = plt.subplots(figsize=(8, 6))
disp.plot(cmap='Blues', xticks_rotation=45, ax=ax)
plt.title("Confusion Matrix - Test Set")
plt.tight_layout()
plt.savefig("confusion_matrix_test_set.png")
plt.show()

# Classification report heatmap
report = classification_report(test_labels, test_preds, output_dict=True)
df_report = pd.DataFrame(report).iloc[:-1, :-1].T
plt.figure(figsize=(8, 5))
sns.heatmap(df_report, annot=True, cmap="YlGnBu")
plt.title("Classification Report Heatmap")
plt.tight_layout()
plt.savefig("classification_report_heatmap.png")
plt.show()

```







```
!pip install matplotlib
```

```
➞ Defaulting to user installation because normal site-packages is not writeable
Collecting matplotlib
  Downloading matplotlib-3.9.4-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (10.2 MB)
Collecting contourpy>=1.0.1 (from matplotlib)
  Downloading contourpy-1.3.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (287 kB)
Collecting cyclor>=0.10 (from matplotlib)
  Downloading cyclor-0.12.1-py3-none-any.whl.metadata (3.8 kB)
Collecting fonttools>=4.22.0 (from matplotlib)
  Downloading fonttools-4.57.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (102.5/102.5 kB) 1.3 MB/s eta 0:00:00
Collecting kiwisolver>=1.3.1 (from matplotlib)
  Downloading kiwisolver-1.4.7-cp39-cp39-manylinux_2_12_x86_64.manylinux2010_x86_64.whl (132 kB)
Requirement already satisfied: numpy>=1.23 in /home/ab12660/.local/lib/python3.9/site-packages (from matplotlib)
Requirement already satisfied: packaging>=20.0 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages (from matplotlib)
Collecting pillow>=8 (from matplotlib)
  Downloading pillow-11.2.1-cp39-cp39-manylinux_2_28_x86_64.whl.metadata (8.9 kB)
Collecting pyparsing>=2.3.1 (from matplotlib)
  Downloading pyparsing-3.2.3-py3-none-any.whl.metadata (5.0 kB)
Requirement already satisfied: python-dateutil>=2.7 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages (from matplotlib)
Collecting importlib-resources>=3.2.0 (from matplotlib)
  Downloading importlib_resources-6.5.2-py3-none-any.whl.metadata (3.9 kB)
Requirement already satisfied: zipp>=3.1.0 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages (from importlib-resources)
Requirement already satisfied: six>=1.5 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages (from pyparsing)
Downloading matplotlib-3.9.4-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (10.2 MB)
  8.3/8.3 MB 48.1 MB/s eta 0:00:00
Downloading contourpy-1.3.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (287 kB)
  321.9/321.9 kB 8.8 MB/s eta 0:00:00
Downloading cyclor-0.12.1-py3-none-any.whl (8.3 kB)
Downloading fonttools-4.57.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (102.5 kB)
  4.6/4.6 MB 53.3 MB/s eta 0:00:00
Downloading importlib_resources-6.5.2-py3-none-any.whl (37 kB)
Downloading kiwisolver-1.4.7-cp39-cp39-manylinux_2_12_x86_64.manylinux2010_x86_64.whl (132 kB)
  1.6/1.6 MB 32.3 MB/s eta 0:00:00
Downloading pillow-11.2.1-cp39-cp39-manylinux_2_28_x86_64.whl (4.6 MB)
  4.6/4.6 MB 46.6 MB/s eta 0:00:00
Downloading pyparsing-3.2.3-py3-none-any.whl (111 kB)
  111.1/111.1 kB 2.8 MB/s eta 0:00:00
Installing collected packages: pyparsing, pillow, kiwisolver, importlib-resources, matplotlib, contourpy, cyclor, fonttools
WARNING: The scripts fonttools, pyftmerge, pyftsubset and ttx are installed to the user base but not in the PATH.
Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
Successfully installed contourpy-1.3.0 cyclor-0.12.1 fonttools-4.57.0 importlib-resources-6.5.2 kiwisolver-1.4.7 matplotlib-3.9.4 pillow-11.2.1 pyparsing-3.2.3
```

```
!pip install seaborn
```

```
➡ Defaulting to user installation because normal site-packages is not writeable
Collecting seaborn
  Downloading seaborn-0.13.2-py3-none-any.whl.metadata (5.4 kB)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in /home/ab12660/.local/lib/python3.9/site-packages (numpy)
Requirement already satisfied: pandas>=1.2 in /home/ab12660/.local/lib/python3.9/site-packages (pandas)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /home/ab12660/.local/lib/python3.9/site-packages (matplotlib)
Requirement already satisfied: contourpy>=1.0.1 in /home/ab12660/.local/lib/python3.9/site-packages (contourpy)
Requirement already satisfied: cycler>=0.10 in /home/ab12660/.local/lib/python3.9/site-packages (cycler)
Requirement already satisfied: fonttools>=4.22.0 in /home/ab12660/.local/lib/python3.9/site-packages (fonttools)
Requirement already satisfied: kiwisolver>=1.3.1 in /home/ab12660/.local/lib/python3.9/site-packages (kiwisolver)
Requirement already satisfied: packaging>=20.0 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages (packaging)
Requirement already satisfied: pillow>=8 in /home/ab12660/.local/lib/python3.9/site-packages (pillow)
Requirement already satisfied: pyparsing>=2.3.1 in /home/ab12660/.local/lib/python3.9/site-packages (pyparsing)
Requirement already satisfied: python-dateutil>=2.7 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages (python-dateutil)
Requirement already satisfied: importlib-resources>=3.2.0 in /home/ab12660/.local/lib/python3.9/site-packages (importlib-resources)
Requirement already satisfied: pytz>=2020.1 in /home/ab12660/.local/lib/python3.9/site-packages (pytz)
Requirement already satisfied: tzdata>=2022.7 in /home/ab12660/.local/lib/python3.9/site-packages (tzdata)
Requirement already satisfied: zipp>=3.1.0 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages (zipp)
Requirement already satisfied: six>=1.5 in /share/apps/pyenv/py3.9/lib/python3.9/site-packages (six)
Downloading seaborn-0.13.2-py3-none-any.whl (294 kB)
----- 294.9/294.9 kB 3.2 MB/s eta 0:00:00
Installing collected packages: seaborn
Successfully installed seaborn-0.13.2
```

```
# Get the logs stored by the Trainer
log_history = peft_lora_finetuning_trainer.state.log_history

# Extract accuracy values per epoch
train_accs = []
eval_accs = []
epochs = []


for entry in log_history:
    if 'eval_accuracy' in entry:
        eval_accs.append(entry['eval_accuracy'])
        epochs.append(entry['epoch'])
    if 'train_accuracy' in entry: # Only if training accuracy is logged
        train_accs.append(entry['train_accuracy']) # Optional; might be missing

# Optional: manually insert training accuracy if it's missing
if not train_accs:
    train_accs = [None] * len(eval_accs)
```


```
import matplotlib.pyplot as plt


plt.figure(figsize=(8, 5))
plt.plot(epochs, eval_accs, label='Validation Accuracy', marker='o', color='tomato')
if any(train_accs):
    plt.plot(epochs, train_accs, label='Training Accuracy', marker='o', color='skyblue')

plt.title("Accuracy Over Epochs")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.ylim(0.5, 1.0)
plt.xticks(epochs)
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.savefig("accuracy_vs_epoch.png") # For your report
plt.show()
```



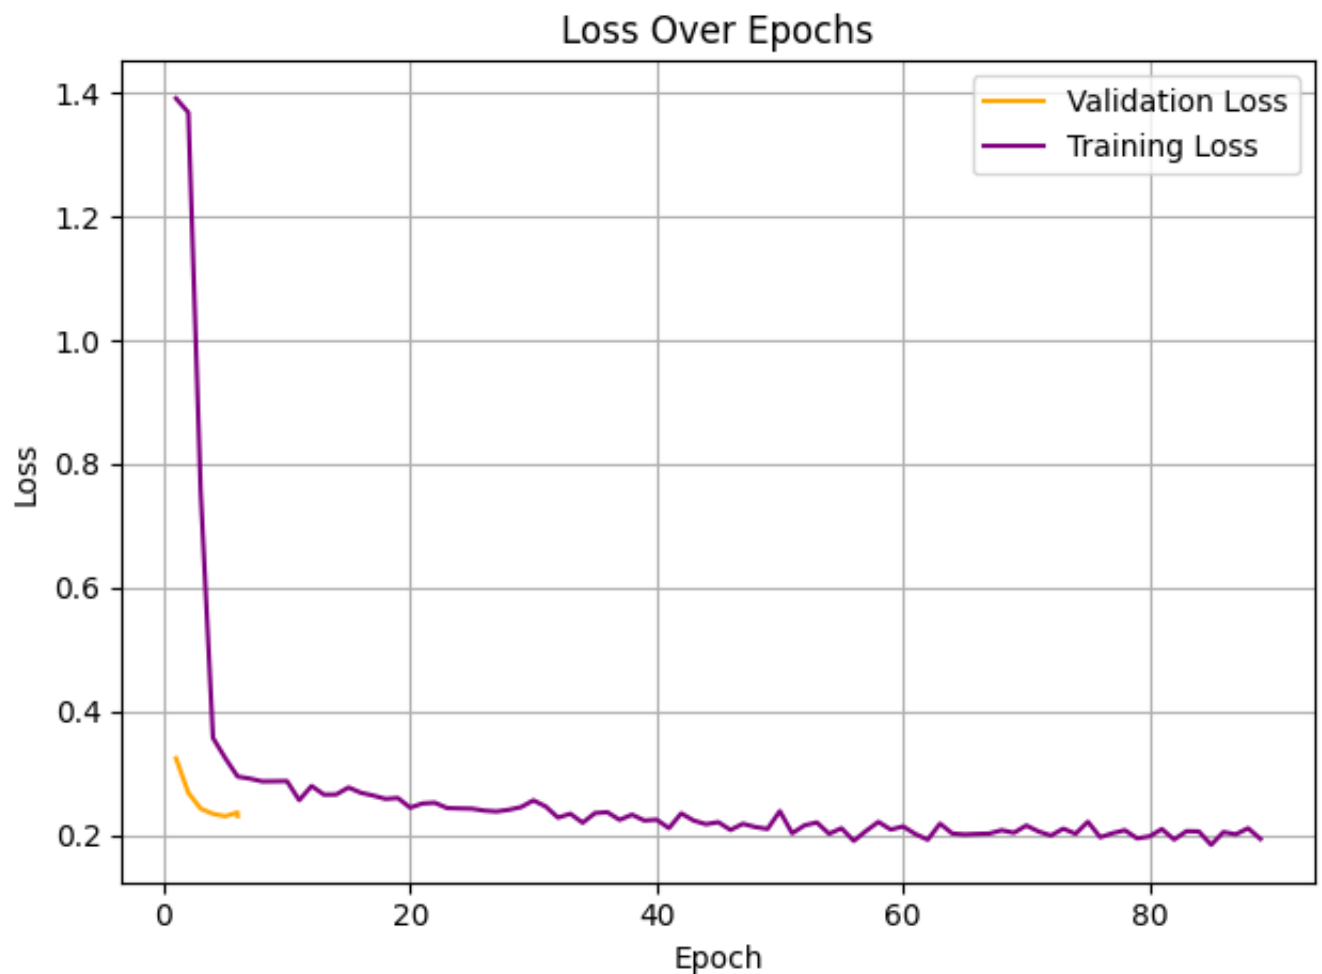
```
plt.plot(epochs, eval_accs, label='Validation Accuracy')
```

 [`<matplotlib.lines.Line2D at 0x14e5651a5bb0>`]





```
train_losses = [entry['loss'] for entry in log_history if 'loss' in entry]
eval_losses = [entry['eval_loss'] for entry in log_history if 'eval_loss' in entry]
epochs_loss = [entry['epoch'] for entry in log_history if 'eval_loss' in entry]

plt.plot(epochs_loss, eval_losses, label='Validation Loss', color='orange')
plt.plot(range(1, len(train_losses)+1), train_losses, label='Training Loss', color='purple')
plt.title("Loss Over Epochs")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig("loss_vs_epoch.png")
plt.show()
```



```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

cm = confusion_matrix(test_labels, test_preds)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
disp.plot(cmap='Blues', xticks_rotation=45)
plt.title("Test Set Confusion Matrix")
plt.tight_layout()
plt.savefig("confusion_matrix.png")
plt.show()
```




```

from sklearn.metrics import classification_report


report = classification_report(test_labels, test_preds, target_names=class_names,
df_report = pd.DataFrame(report).transpose()
df_report.to_csv("classification_report.csv")

# Optional visualization:
import seaborn as sns
plt.figure(figsize=(8, 5))
sns.heatmap(df_report.iloc[:-1, :-1], annot=True, cmap="YlGnBu", fmt=".2f")
plt.title("Classification Report Heatmap")
plt.tight_layout()
plt.savefig("classification_report_heatmap.png")
plt.show()

```



```
plt.bar(['Train Accuracy', 'Test Accuracy'], [train_acc, test_acc], color=['skyblu  
plt.ylim(0, 1)  
plt.ylabel('Accuracy')  
plt.title('Train vs Test Accuracy')  
plt.grid(axis='y')  
plt.tight_layout()  
plt.savefig("train_vs_test_accuracy.png")  
plt.show()
```



```
examples = [  
    "NASA launches a new satellite into orbit.",  
    "Stocks rally as Fed pauses interest rate hikes.",  
    "Manchester United defeats Liverpool in derby match.",  
    "New AI model breaks records in image generation."  
]
```

```
for sentence in examples:  
    pred = classify(peft_model, tokenizer, sentence)  
    print(f"Text: {sentence}\n → Predicted: {pred}\n")
```



```
Class: 3, Label: Sci/Tech, Text: NASA launches a new satellite into orbit.  
Text: NASA launches a new satellite into orbit.  
→ Predicted: Sci/Tech
```

```
Class: 2, Label: Business, Text: Stocks rally as Fed pauses interest rate hikes.  
Text: Stocks rally as Fed pauses interest rate hikes.  
→ Predicted: Business
```

```
Class: 1, Label: Sports, Text: Manchester United defeats Liverpool in derby match.  
Text: Manchester United defeats Liverpool in derby match.  
→ Predicted: Sports
```

```
Class: 3, Label: Sci/Tech, Text: New AI model breaks records in image generation.  
Text: New AI model breaks records in image generation.  
→ Predicted: Sci/Tech
```

Start coding or generate with AI.

